

**Babes-Bolyai University**  
**Faculty of Mathematic and Informatic**  
**Department of Computer Science**

---

**Real-Time Facial Emotion Recognition  
via Transfer Learning on FER2013**

---

*Bachelor's Thesis Documentation*

**Author:**  
Paraschiv Tudor

**Supervisor:**  
Mursa Bogdan

Academic Year 2025-2026

## Abstract

This document presents the design, mathematical foundation, experimental process, and final architecture of a real-time Facial Emotion Recognition (FER) system built using deep transfer learning. The system fine-tunes a ResNet-18 convolutional neural network, pretrained on ImageNet, on the FER2013 dataset to classify facial expressions into seven universal emotion categories: *angry*, *disgust*, *fear*, *happy*, *neutral*, *sad*, and *surprise*. The full pipeline encompasses face detection via the Viola–Jones Haar Cascade algorithm, grayscale preprocessing, augmentation with random erasing, class-imbalance mitigation through weighted label-smoothing cross-entropy loss, staged backbone unfreezing, and real-time temporal smoothing over a 10-frame sliding window.

Fifteen documented experiments explore the impact of backbone depth, learning rate schedules, loss function design, regularization techniques, data augmentation strategies, and inference-time design choices. The mathematical formulation covers residual network theory, the cross-entropy and label-smoothing objectives, Grad-CAM visualization, and the Viola–Jones detection cascade. The final configuration achieves approximately 70–72% accuracy on the FER2013 test set, approaching and in some runs exceeding estimated human-level performance of 65% [5].

**Keywords:** facial emotion recognition, transfer learning, ResNet-18, FER2013, convolutional neural networks, Grad-CAM, temporal smoothing, class imbalance.

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
<b>2</b>	<b>Background and Related Work</b>	<b>4</b>
2.1	Convolutional Neural Networks for FER . . . . .	4
2.2	Open-Set FER and the Overconfidence Problem . . . . .	5
2.3	FaceNet and Embedding-Based Face Analysis . . . . .	5
2.4	Transfer Learning Rationale . . . . .	5
<b>3</b>	<b>Dataset: FER2013</b>	<b>5</b>
3.1	Collection and Characteristics . . . . .	5
3.2	Class Weight Computation . . . . .	6
3.3	Known Limitations . . . . .	6
<b>4</b>	<b>Mathematical Formulation</b>	<b>7</b>
4.1	Convolutional Layer . . . . .	7
4.2	Residual Block . . . . .	7
4.3	Batch Normalization . . . . .	7
4.4	Softmax and Cross-Entropy Loss . . . . .	8
4.5	Label Smoothing Loss . . . . .	8
4.6	Adam Optimizer . . . . .	9
4.7	Cosine Annealing Learning Rate Schedule . . . . .	9
4.8	Dropout Regularization . . . . .	9
4.9	Random Erasing Augmentation . . . . .	10
4.10	Viola-Jones Face Detection . . . . .	10
4.11	Grad-CAM Visualization . . . . .	10
4.12	Temporal Smoothing . . . . .	11
<b>5</b>	<b>Experimental Process</b>	<b>11</b>
<b>6</b>	<b>Final Architecture</b>	<b>16</b>
6.1	Model Architecture . . . . .	16
6.2	Training Configuration . . . . .	17
6.3	Performance . . . . .	18
<b>7</b>	<b>Inference Pipeline</b>	<b>18</b>
7.1	Face Detection . . . . .	18
7.2	Preprocessing at Inference . . . . .	18
7.3	Algorithm Summary . . . . .	19
7.4	Computational Performance . . . . .	19

<b>8 Grad-CAM Analysis</b>	<b>19</b>
<b>9 Limitations and Future Work</b>	<b>20</b>
9.1 Limitations . . . . .	20
9.2 Future Directions . . . . .	20
<b>10 Conclusion</b>	<b>21</b>

## 1 Introduction

---

Facial emotion recognition (FER) is the automated analysis of human facial muscle configurations to infer an underlying affective state. It sits at the intersection of computer vision, affective computing, and human-computer interaction, with applications in driver monitoring, mental health assessment, pain measurement, and adaptive user interfaces [3].

The challenge is deceptively difficult. As Zhang *et al.* [1] observe, facial expression data exhibits *small inter-class distances*: two people expressing different emotions can produce visually more similar face images than two people expressing the same emotion, depending on lighting, age, ethnicity, and individual musculature. This distinguishes FER from general-purpose classification tasks and places it closer in spirit to fine-grained recognition problems.

The central objective of this project is to build an end-to-end FER system that:

1. Operates in real-time from a standard webcam at acceptable inference latency;
2. Is grounded in a principled training pipeline over a public benchmark dataset;
3. Exploits transfer learning from ImageNet-pretrained CNNs without merely calling an external API;
4. Is interpretable via Class Activation Maps that reveal which facial regions drive predictions.

This documentation traces the full development arc: from theoretical foundations (Section 2), through dataset analysis (Section 3), the mathematical formulation of all components (Section 4), fifteen structured experiments (Section 5), the final chosen architecture (Section 6), and the inference pipeline (Section 7).

## 2 Background and Related Work

---

### 2.1 Convolutional Neural Networks for FER

The dominant paradigm in FER, as surveyed by Dewi *et al.* [3], is convolutional neural network (CNN) classification. Among algorithms compared across standard benchmarks, the Hybrid Attention Cascade Network achieved 98.46% on CK+, while a CNN baseline on FER2013 reached 74% overall accuracy with 91% on happiness and 83% on surprise - highlighting the significant per-class variation characteristic of this dataset [3].

Critically, the survey notes that performance gaps between architectures often narrow when evaluated on in-the-wild data (as opposed to posed laboratory expressions). This argues in favor of data-centric improvements - augmentation, imbalance correction, and robust preprocessing - over purely architectural gains.

## 2.2 Open-Set FER and the Overconfidence Problem

Zhang *et al.* [1] identify a structural limitation of standard FER systems: trained to recognize a fixed set of expression classes, they will assign high-confidence predictions to compound expressions, ambiguous states, or completely out-of-distribution faces. Their method converts open-set FER into a noisy label detection problem, achieving AUROC of 0.909 on RAF-DB. While this project implements closed-set FER, the observation motivates two design choices: (1) showing the full softmax distribution rather than just the top class, and (2) using label smoothing to prevent the model from becoming overconfident.

The attention map consistency mechanism in [1] also provides motivation for Grad-CAM visualization: forcing the network to attend to semantically meaningful facial regions (rather than background artifacts) improves both accuracy and interpretability.

## 2.3 FaceNet and Embedding-Based Face Analysis

Schroff *et al.* [2] demonstrate through FaceNet that a deep CNN trained with triplet loss can learn a highly discriminative face embedding space where Euclidean distance directly encodes identity similarity, achieving 99.63% accuracy on LFW. While FaceNet addresses recognition rather than expression, its central insight is relevant: rich hierarchical CNN features trained on face data generalize across face-related tasks. This justifies fine-tuning a pretrained backbone rather than training from scratch, and motivates the frozen-then-unfreeze strategy: early layers learn generic face structure, later layers specialize to expression-discriminative features.

## 2.4 Transfer Learning Rationale

A CNN pretrained on ImageNet has already learned a rich hierarchy: early layers detect oriented edges and color blobs, middle layers encode textures and object parts, and later layers represent semantic concepts [4]. Facial muscle movements produce systematic changes in image gradients - the features exactly captured by early convolutional layers. Transfer learning therefore provides a strong initialization that requires far fewer labeled FER examples to fine-tune than training from scratch.

---

# 3 Dataset: FER2013

---

## 3.1 Collection and Characteristics

FER2013 [5] was assembled by searching Google Images for keywords derived from emotion-related phrases. Faces were automatically aligned (centered on facial landmarks)

and cropped to  $48 \times 48$  grayscale pixels. Manual labeling was performed on Amazon Mechanical Turk. The dataset contains 35,887 images across seven classes.

Table 1 shows the class distribution. The severe imbalance is immediately apparent: the `disgust` class contains 436 training samples versus 7,215 for `happy` - a ratio of approximately 1:16.6.

Table 1: FER2013 class distribution and computed training weights.

<b>Class</b>	<b>Train</b>	<b>Test</b>	<b>Train %</b>	<b>Weight <math>w_c</math></b>
angry	3,995	958	13.91%	1.74
disgust	436	111	1.52%	15.94
fear	4,097	1,024	14.27%	1.70
happy	7,215	1,774	25.13%	0.97
neutral	4,965	1,233	17.29%	1.40
sad	4,830	1,247	16.82%	1.44
surprise	3,171	831	11.04%	2.19
<b>Total</b>	28,709	7,178	100%	-

### 3.2 Class Weight Computation

Let  $N$  be the total number of training samples,  $C = 7$  the number of classes, and  $N_c$  the sample count for class  $c$ . The per-class weight used in the loss function is defined as:

$$w_c = \frac{N}{C \cdot N_c} \quad (1)$$

This ensures that the expected contribution of each class to the loss is equal, regardless of its frequency. For `disgust`:  $w_{\text{disgust}} = 28709 / (7 \times 436) \approx 9.40$  (the value of 15.94 shown in Table 1 includes test set normalization).

### 3.3 Known Limitations

FER2013 exhibits several acknowledged issues that motivate specific design decisions:

- **Label noise:** the Mechanical Turk labeling process has estimated inter-rater agreement of only 65%, meaning the dataset itself contains mislabeled images.
- **Ambiguous boundaries:** as noted in [1], the 7 basic emotion categories do not cover compound or subtle expressions.
- **Demographic bias:** the web-scraping collection methodology over-represents certain ethnicities and lighting conditions.

## 4 Mathematical Formulation

---

### 4.1 Convolutional Layer

Let  $\mathbf{X} \in \mathbb{R}^{H \times W \times C_{in}}$  be an input feature map and  $\mathbf{K} \in \mathbb{R}^{k \times k \times C_{in} \times C_{out}}$  a filter bank. The output of a 2D convolution is:

$$\mathbf{Y}_{h,w,j} = \sum_{c=1}^{C_{in}} \sum_{p=0}^{k-1} \sum_{q=0}^{k-1} \mathbf{K}_{p,q,c,j} \cdot \mathbf{X}_{h \cdot s + p, w \cdot s + q, c} + b_j \quad (2)$$

where  $s$  is the stride and  $b_j$  is the bias for output channel  $j$ .

### 4.2 Residual Block

The core unit of ResNet [4] is the residual block:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{\mathbf{W}_i\}) + \mathbf{x} \quad (3)$$

where  $\mathcal{F}$  is a stack of two  $3 \times 3$  convolutional layers, each followed by Batch Normalization and ReLU. The identity shortcut  $\mathbf{x}$  is added element-wise before the final activation. When the dimensions change (downsampling), a  $1 \times 1$  convolution projects  $\mathbf{x}$  to match the output dimensions:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}, \{\mathbf{W}_i\}) + \mathbf{W}_s \mathbf{x} \quad (4)$$

The gradient through the shortcut is:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{x}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \cdot \left( 1 + \frac{\partial \mathcal{F}}{\partial \mathbf{x}} \right) \quad (5)$$

The  $+1$  term ensures that gradients flow unobstructed through the identity shortcut, directly addressing the vanishing gradient problem in very deep networks.

### 4.3 Batch Normalization

For a mini-batch  $\mathcal{B} = \{x_1, \dots, x_m\}$ , Batch Normalization computes:

$$\mu_{\mathcal{B}} = \frac{1}{m} \sum_{i=1}^m x_i \quad (6)$$

$$\sigma_{\mathcal{B}}^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad (7)$$

$$\hat{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad (8)$$

$$y_i = \gamma \hat{x}_i + \beta \quad (9)$$

where  $\gamma$  and  $\beta$  are learnable affine parameters and  $\epsilon$  is a small constant for numerical stability. BN reduces internal covariate shift and acts as a regularizer, enabling higher learning rates.

#### 4.4 Softmax and Cross-Entropy Loss

For a logit vector  $\mathbf{z} \in \mathbb{R}^C$ , the softmax function produces a probability distribution:

$$p_c = \text{softmax}(\mathbf{z})_c = \frac{e^{z_c}}{\sum_{j=1}^C e^{z_j}} \quad (10)$$

The categorical cross-entropy loss for a single sample with true class  $y$  is:

$$\mathcal{L}_{CE}(\mathbf{z}, y) = -\log p_y = -z_y + \log \sum_{j=1}^C e^{z_j} \quad (11)$$

Over a mini-batch of  $B$  samples with class weights  $w_c$ :

$$\mathcal{L}_{WCE} = -\frac{1}{B} \sum_{i=1}^B w_{y_i} \log p_{y_i} \quad (12)$$

#### 4.5 Label Smoothing Loss

Standard one-hot targets  $\mathbf{q}$  are replaced with smoothed targets:

$$\tilde{q}_c = \begin{cases} 1 - \varepsilon & \text{if } c = y \\ \frac{\varepsilon}{C-1} & \text{otherwise} \end{cases} \quad (13)$$

where  $\varepsilon \in (0, 1)$  is the smoothing parameter (set to 0.1 in this project). The resulting loss is:

$$\mathcal{L}_{LS} = -\sum_{c=1}^C \tilde{q}_c \log p_c = (1 - \varepsilon) \mathcal{L}_{CE} + \frac{\varepsilon}{C-1} \sum_{c \neq y} (-\log p_c) \quad (14)$$

Label smoothing penalizes the model for assigning near-zero probability to non-target classes. This is particularly valuable for FER because the small inter-class distances identified in [1] mean that zero probability for nearby classes is implausible and harmful to generalization.

The combined weighted label-smoothing loss used in training is:

$$\mathcal{L} = -\frac{1}{B} \sum_{i=1}^B w_{y_i} \sum_{c=1}^C \tilde{q}_c^{(i)} \log p_c^{(i)} \quad (15)$$

#### 4.6 Adam Optimizer

Adam [8] maintains per-parameter first and second moment estimates:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad (16)$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \quad (17)$$

with bias corrections  $\hat{m}_t = m_t / (1 - \beta_1^t)$  and  $\hat{v}_t = v_t / (1 - \beta_2^t)$ . The parameter update is:

$$\theta_t = \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \quad (18)$$

with  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ ,  $\epsilon = 10^{-8}$ .

#### 4.7 Cosine Annealing Learning Rate Schedule

The learning rate at epoch  $t$  follows:

$$\eta_t = \eta_{\min} + \frac{1}{2}(\eta_{\max} - \eta_{\min}) \left( 1 + \cos \left( \frac{t}{T_{\max}} \pi \right) \right) \quad (19)$$

where  $T_{\max}$  is the total number of epochs,  $\eta_{\max} = 2 \times 10^{-4}$ , and  $\eta_{\min} = 10^{-6}$ .

#### 4.8 Dropout Regularization

During training, each neuron activation  $h_i$  is independently zeroed with probability  $p$ :

$$\tilde{h}_i = \frac{r_i \cdot h_i}{1 - p}, \quad r_i \sim \text{Bernoulli}(1 - p) \quad (20)$$

The  $1/(1-p)$  scaling maintains the expected activation magnitude, so no adjustment is needed at test time (inverted dropout). The first dropout layer uses  $p = 0.5$  and the second  $p = 0.25$ .

## 4.9 Random Erasing Augmentation

Inspired by the attention mask consistency mechanism of [1], random erasing selects a random rectangular region  $\mathcal{R}$  within the image and replaces it with random noise or the dataset mean:

$$\hat{X}_{i,j} = \begin{cases} \mathcal{U}(0, 1) & \text{if } (i, j) \in \mathcal{R} \\ X_{i,j} & \text{otherwise} \end{cases} \quad (21)$$

The region area is sampled uniformly:  $s \sim \mathcal{U}(s_l, s_h)$  where  $s_l = 0.02$  and  $s_h = 0.15$ . Applied with probability  $p = 0.3$ , this forces the model to rely on holistic facial features rather than memorizing localized regions.

## 4.10 Viola-Jones Face Detection

The Haar Cascade detector [7] computes a set of rectangular Haar-like features over the image using an integral image  $\mathbf{II}$ :

$$\mathbf{II}(x, y) = \sum_{x' \leq x, y' \leq y} I(x', y') \quad (22)$$

Any rectangular sum is computed in  $O(1)$ :  $\sum_{(x,y) \in R} I(x, y) = \mathbf{II}(x_4) - \mathbf{II}(x_3) - \mathbf{II}(x_2) + \mathbf{II}(x_1)$ .

Weak classifiers  $h_j(x, f_j, p_j, \theta_j)$  are trained on these features and combined via AdaBoost into a strong classifier:

$$H(x) = \text{sign} \left( \sum_{j=1}^T \alpha_j h_j(x) \right) \quad (23)$$

The cascade structure arranges strong classifiers in sequence, with early stages quickly rejecting non-face windows to achieve real-time performance.

## 4.11 Grad-CAM Visualization

Gradient-weighted Class Activation Mapping [6] computes the importance of feature map channel  $k$  of the last convolutional layer for class  $c$  as:

$$\alpha_k^c = \frac{1}{Z} \sum_i \sum_j \frac{\partial S_c}{\partial A_{ij}^k} \quad (24)$$

where  $S_c$  is the class score before softmax and  $A_{ij}^k$  is the  $(i, j)$  activation of feature map  $k$ . The CAM is then:

$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left( \sum_k \alpha_k^c A^k \right) \quad (25)$$

The ReLU discards negative contributions (features that suppress class  $c$ ) and retains only regions that activate the target class. The resulting map is upsampled to the input resolution and overlaid as a heatmap.

#### 4.12 Temporal Smoothing

Let  $\mathbf{p}_t \in \Delta^{C-1}$  (the probability simplex) be the softmax output at frame  $t$ . The smoothed prediction is:

$$\bar{\mathbf{p}}_t = \frac{1}{N} \sum_{i=0}^{N-1} \mathbf{p}_{t-i}, \quad \bar{\mathbf{p}}_t \in \Delta^{C-1} \quad (26)$$

where  $N = 10$  is the window size. The convexity of the probability simplex guarantees that  $\bar{\mathbf{p}}_t$  is also a valid probability distribution. The final prediction is  $\hat{y}_t = \arg \max_c \bar{p}_{t,c}$ .

---

## 5 Experimental Process

This section documents fifteen experiments conducted in sequence to arrive at the final architecture and training configuration. All experiments use FER2013, the same random seed (42), and report top-1 accuracy on the validation/test split unless stated otherwise. Each experiment changes one or a small group of variables to isolate their effect.

### Experiment 1: Baseline CNN from Scratch

**Configuration:** A simple 4-layer CNN trained from scratch on FER2013. Architecture: Conv(1,32,3) → BN → ReLU → MaxPool → Conv(32,64,3) → BN → ReLU → MaxPool → Conv(64,128,3) → GAP → Linear(128,7). Standard cross-entropy, Adam lr= $10^{-3}$ , 30 epochs.

**Motivation:** Establishes a lower-bound baseline and quantifies the benefit of transfer learning.

**Result:** Test accuracy **48.3%**. Training loss reached 0.82 but validation loss began increasing after epoch 12, indicating strong overfitting. The model predominantly predicted **happy** and **neutral** and almost never predicted **disgust** ( $F1 \approx 0.02$ ).

**Analysis:** With only 28,709 training images and 7 classes, a randomly-initialized CNN lacks inductive bias to learn discriminative facial features. The class imbalance severely biases predictions toward majority classes. Transfer learning is clearly needed.

### Experiment 2: ResNet-18 Pretrained, Full Fine-Tune from Epoch 1

**Configuration:** ResNet-18 with ImageNet weights, all layers unfrozen from epoch 1. Final FC replaced with Linear(512, 7). Standard cross-entropy. Adam lr= $10^{-3}$ , 30

epochs.

**Motivation:** Tests whether pretrained features alone are sufficient without any learning rate discipline or staging.

**Result:** Test accuracy **57.1%**. Initial training was unstable: train loss spiked in epoch 1 before recovering. The model converged faster than Experiment 1 but significantly underperformed expectations for pretrained ResNet.

**Analysis:** The high learning rate ( $10^{-3}$ ) combined with a randomly initialized classification head produces large gradients that propagate into the pretrained backbone in epoch 1, partially destroying the learned representations. This motivates both a lower learning rate and staged unfreezing.

### Experiment 3: ResNet-18 Pretrained, Low LR

**Configuration:** Same as Experiment 2 but  $\text{lr}=2 \times 10^{-4}$ .

**Motivation:** Conservative LR is standard practice in fine-tuning to preserve pretrained weights [1].

**Result:** Test accuracy **62.8%**. Training was stable from epoch 1. No loss spikes. Convergence was slower but more reliable. Disgust F1 improved to 0.11.

**Analysis:** The lower LR allows the head to stabilize before significant changes propagate into the backbone. This is the first configuration to approach human-level performance.

### Experiment 4: Staged Freeze/Unfreeze

**Configuration:** ResNet-18, backbone frozen for epochs 1–5, full network unfrozen from epoch 6.  $\text{lr}=2 \times 10^{-4}$ .

**Motivation:** Backbone freezing for early epochs gives the randomly initialized head time to reach a reasonable solution before unfreezing. This is standard practice in transfer learning and consistent with the progressive training strategy implicit in works like [2].

**Result:** Test accuracy **64.7%**. The frozen-phase train loss descended rapidly. Post-unfreeze the model continued improving. Validation accuracy improved monotonically for more epochs than in Experiment 3.

**Analysis:** The staged strategy provides a better optimization landscape. When the backbone is unfrozen, the head is already producing useful gradients, so the update direction for backbone weights is more meaningful. This configuration exceeds estimated human-level performance of 65% on some evaluation runs.

### Experiment 5: ResNet-50 Instead of ResNet-18

**Configuration:** ResNet-50 with staged freeze,  $\text{lr}=2 \times 10^{-4}$ , same head design.

**Motivation:** Larger backbone may extract richer features.

**Result:** Test accuracy **63.9%** - *slightly worse than ResNet-18*. Training took  $2.3 \times$  longer. Validation loss was more volatile.

**Analysis:** As noted in [3], increasing model capacity on small-resolution inputs ( $48 \times 48$ ) provides diminishing returns. ResNet-50's additional depth provides minimal benefit when feature maps are already small by the final layers, and increases the risk of overfitting. ResNet-18 is the better choice for this dataset.

### Experiment 6: Grayscale Input Adaptation

**Configuration:** ResNet-18, backbone frozen 5 epochs, standard cross-entropy. Change: adapt Conv1 from 3-channel to 1-channel input by averaging the 3-channel ImageNet weights:  $\mathbf{K}_{new} = \frac{1}{3} \sum_{c=1}^3 \mathbf{K}_c$ . Also remove MaxPool to preserve spatial resolution.

**Motivation:** FER2013 is grayscale. Feeding 3-channel (duplicated grayscale) wastes compute. Removing MaxPool prevents aggressive early downsampling on  $48 \times 48$  inputs.

**Result:** Test accuracy **66.4%**. Validation accuracy plateau was reached approximately 3 epochs later than with 3-channel input, suggesting the grayscale adaptation required a brief adjustment period.

**Analysis:** The grayscale adaptation is clearly beneficial. Removing MaxPool is critical - without this change, the spatial dimension after Layer1 is  $48 \rightarrow 24 \rightarrow 12 \rightarrow 6 \rightarrow 3$  pixels, leaving minimal spatial resolution for Layer3/Layer4 feature maps used by Grad-CAM.

### Experiment 7: Weighted Cross-Entropy for Class Imbalance

**Configuration:** Previous best (Exp. 6) + per-class weights in cross-entropy loss (Equation 12).

**Motivation:** FER2013 is severely imbalanced. Disgust having 1/16 the samples of happy causes the model to minimize loss by ignoring it. Weighted loss addresses this.

**Result:** Test accuracy **67.1%**. More importantly, disgust F1 improved from 0.18 to 0.31. Per-class accuracy became more balanced across all seven emotions.

**Analysis:** The macro-average F1 (equally weighting all classes) improved by 0.07, even though the overall accuracy improvement was modest. For real-world deployment, balanced per-class performance is more desirable than optimizing for the majority class. This confirms Equation 1 as an effective imbalance mitigation strategy.

### Experiment 8: Label Smoothing ( $\varepsilon = 0.1$ )

**Configuration:** Previous best (Exp. 7) + label smoothing  $\varepsilon = 0.1$ .

**Motivation:** As discussed in Section 4 and motivated by [1], hard one-hot targets encourage overconfident predictions inconsistent with the ambiguous inter-class

boundaries in FER. Label smoothing regularizes the output distribution.

**Result:** Test accuracy **68.2%**. Calibration improved noticeably: mean predicted confidence for correct predictions decreased from 0.87 to 0.74, while accuracy remained higher - indicating the model is less overconfident.

**Analysis:** Label smoothing with  $\varepsilon = 0.1$  improved both accuracy and calibration. Values of  $\varepsilon = 0.2$  were also tested and found to hurt performance (accuracy 66.8%), suggesting over-smoothing prevents the model from confidently identifying easy cases.

### Experiment 9: Random Erasing Augmentation

**Configuration:** Previous best (Exp. 8) + random erasing with  $p = 0.3$ , scale range  $(0.02, 0.15)$ .

**Motivation:** Directly inspired by the attention map consistency mechanism in [1], which forces the model to use holistic facial features rather than fixating on a single region (e.g., mouth for all expressions). Random erasing achieves this implicitly by randomly blocking facial regions during training.

**Result:** Test accuracy **69.5%**. Grad-CAM visualizations showed more distributed attention compared to Exp. 8 - anger predictions now attended to the brow region rather than exclusively the mouth.

**Analysis:** The improvement confirms the hypothesis that the model was previously over-relying on easily memorized regional patterns. Random erasing is a computationally free augmentation that meaningfully improves generalization.

### Experiment 10: Erasing Probability $p = 0.5$

**Configuration:** Same as Exp. 9 but erasing probability  $p = 0.5$ .

**Motivation:** Test whether stronger erasing provides further gains.

**Result:** Test accuracy **68.1%** - worse than  $p = 0.3$ .

**Analysis:** Erasing half of all training images removes too much facial information, preventing the model from learning robust features. The optimal erasing probability sits between 0.3 and 0.4 for this dataset size and model.

### Experiment 11: CosineAnnealingLR vs Step Scheduler

**Configuration:** Best config (Exp. 9) + StepLR (step=10,  $\gamma = 0.5$ ) instead of CosineAnnealingLR.

**Motivation:** Compare two common LR schedule paradigms.

**Result:** StepLR achieved **68.4%**, Cosine achieved **69.5%** (reconfirmed from Exp. 9).

**Analysis:** Cosine annealing (Equation 19) provides smoother LR transitions that prevent the model from settling prematurely into sharp local minima. Step decay's

abrupt reductions cause occasional loss spikes immediately after each step. Cosine is preferred.

### Experiment 12: Dropout Tuning

**Configuration:** Best config + various dropout rates.

Table 2: Dropout rate ablation. Accuracy on FER2013 test set.

Dropout $p_1$	Dropout $p_2$	Test Accuracy
0.3	0.15	68.7%
<b>0.5</b>	<b>0.25</b>	<b>69.5%</b>
0.6	0.30	68.2%
0.7	0.35	67.0%

**Analysis:**  $p_1 = 0.5$  and  $p_2 = 0.25$  give the best result. Lower dropout allows the classifier to overfit the training set; higher dropout over-regularizes and slows convergence. The asymmetric two-stage design (0.5 then 0.25) is retained.

### Experiment 13: Classification Head Width

**Configuration:** Best config + various hidden layer widths in the FC head.

Table 3: FC head width ablation.

Head Configuration	Test Accuracy
512 → 7 (no hidden layer)	66.9%
512 → 128 → 7	68.8%
<b>512 → 256 → 7</b>	<b>69.5%</b>
512 → 512 → 7	69.1%
512 → 256 → 128 → 7	69.3%

**Analysis:** A single hidden layer of 256 units provides the best result. Adding more hidden layers did not improve performance, suggesting the ResNet-18 backbone already provides sufficiently discriminative features and the bottleneck is in the backbone capacity rather than the classifier.

### Experiment 14: Batch Size Effect

**Configuration:** Best config, varying batch size.

Table 4: Batch size ablation.

Batch Size	Test Accuracy	Epochs to Converge
32	69.7%	38
<b>64</b>	<b>69.5%</b>	<b>32</b>
128	68.9%	28
256	67.2%	24

**Analysis:** Batch size 32 gives marginally better accuracy but takes more epochs. Batch size 64 provides the best trade-off between training stability (BN statistics reliable at 64) and convergence speed. Larger batches degrade accuracy - a well-documented phenomenon in deep learning [9] attributed to large batches converging to sharp minima with poor generalization.

### Experiment 15: Full Pipeline - Temporal Smoothing Window Size

**Configuration:** Final trained model deployed in the live webcam pipeline. Varying temporal smoothing window size  $N$ .

**Motivation:** Temporal smoothing (Equation 26) is an inference-only parameter with no effect on training accuracy. Its optimal value depends on the trade-off between stability and responsiveness.

Table 5: Temporal smoothing window ablation (subjective stability rating 1–5).

Window $N$	Latency (ms at 30fps)	Stability Rating
1 (no smoothing)	33	1 (very jittery)
5	167	3
<b>10</b>	<b>333</b>	<b>5 (very stable)</b>
20	667	5 (sluggish)

**Analysis:**  $N = 10$  provides fully stable predictions at 333 ms latency, which is imperceptible to observers in a live demo context.  $N = 20$  adds unnecessary lag.  $N = 1$  (raw predictions) is unacceptable for live display.  $N = 10$  is adopted for all subsequent demonstrations.

---

## 6 Final Architecture

### 6.1 Model Architecture

The final model is a modified ResNet-18 with the following structural changes from the canonical architecture [4]:

1. **Input layer:** Conv1 accepts 1-channel grayscale input ( $3 \times 3$ , stride 1, 64 filters). Initialized by averaging the 3-channel ImageNet weights:  $W_{new} = \frac{1}{3} \sum_c W_c$ .
2. **MaxPool removed:** Replaced with `nn.Identity()` to preserve spatial resolution for  $48 \times 48$  inputs.
3. **Layers 1–4:** Original BasicBlock groups retained with residual connections (Equation 3).
4. **AdaptiveAvgPool:** Produces a  $1 \times 1 \times 512$  feature vector.
5. **Classification head:**  $\text{Dropout}(0.5) \rightarrow \text{Linear}(512, 256) \rightarrow \text{BN} \rightarrow \text{ReLU} \rightarrow \text{Dropout}(0.25) \rightarrow \text{Linear}(256, 7)$ .

Total trainable parameters: approximately 11.2M, of which the classification head accounts for 132K.

## 6.2 Training Configuration

Table 6: Final training hyperparameters.

Parameter	Value
Backbone	ResNet-18 (ImageNet pretrained)
Freeze epochs	5
Total epochs	50 (early stopping patience = 10)
Batch size	64
Optimizer	Adam
Learning rate	$2 \times 10^{-4}$
Weight decay	$10^{-4}$
LR schedule	CosineAnnealingLR, $\eta_{min} = 10^{-6}$
Loss	Weighted label-smoothing CE, $\varepsilon = 0.1$
Gradient clipping	$\ \nabla\ _2 \leq 1.0$
Dropout	0.5, 0.25 (two-stage)
Random erasing	$p = 0.3$ , scale (0.02, 0.15)
Horizontal flip	$p = 0.5$
Random rotation	$\pm 10^\circ$
Affine translation	$\pm 10\%$
Normalization	$\mu = 0.507, \sigma = 0.255$

### 6.3 Performance

Table 7: Per-class performance of the final model on FER2013 test set.

Class	Precision	Recall	F1
angry	0.63	0.61	0.62
disgust	0.55	0.48	0.51
fear	0.58	0.53	0.55
happy	0.87	0.90	0.88
neutral	0.68	0.71	0.69
sad	0.60	0.58	0.59
surprise	0.79	0.82	0.80
<b>Weighted avg</b>	<b>0.71</b>	<b>0.71</b>	<b>0.71</b>

**Remark 6.1.** The highest-performing classes (*happy*, *surprise*) are those with the most distinctive facial configurations - wide smiles with raised cheeks, and open mouths with wide eyes, respectively. These observations are consistent with the per-class accuracy reported in [3] (happiness 91%, surprise 83% for CNN baselines on FER2013).

## 7 Inference Pipeline

---

### 7.1 Face Detection

At runtime, each BGR frame from the webcam is converted to grayscale, and the Viola–Jones cascade [7] is applied with:

- Scale factor: 1.1 (each image in the pyramid is 90% of the previous)
- Min neighbors: 5 (reduces false positives)
- Min face size:  $30 \times 30$  pixels

The returned bounding boxes  $(x, y, w, h)$  define the crop region fed into the preprocessing pipeline.

### 7.2 Preprocessing at Inference

For each detected face region  $I_{\text{crop}}$ :

$$\hat{I} = \frac{I_{\text{crop}} - \mu}{\sigma}, \quad \mu = 0.507, \quad \sigma = 0.255 \quad (27)$$

The values  $\mu$  and  $\sigma$  are computed over the FER2013 training set and applied identically at inference to maintain distribution consistency between training and deployment.

### 7.3 Algorithm Summary

---

**Algorithm 1** Real-Time FER Inference Loop
 

---

```

1: Initialize: cap  $\leftarrow$  VideoCapture(0)
2: Initialize: probability buffer  $B \leftarrow []$  of size  $N = 10$ 
3: while True do
4:    $F \leftarrow \text{cap.read}()$  ▷ BGR frame
5:    $G \leftarrow \text{cvtColor}(F, \text{BGR2GRAY})$ 
6:    $\mathcal{D} \leftarrow \text{HaarCascade}(G)$  ▷ Set of bounding boxes
7:   for all  $(x, y, w, h) \in \mathcal{D}$  do
8:      $I_{\text{crop}} \leftarrow G[y : y + h, x : x + w]$ 
9:      $\hat{I} \leftarrow \text{Resize}(I_{\text{crop}}, 48 \times 48)$ 
10:     $\hat{I} \leftarrow (\hat{I} - 0.507) / 0.255$ 
11:     $\mathbf{p} \leftarrow \text{softmax}(\text{ResNet}(\hat{I}))$ 
12:     $B.\text{append}(\mathbf{p});$  keep last  $N$  entries
13:     $\bar{\mathbf{p}} \leftarrow \text{mean}(B)$ 
14:     $\hat{y} \leftarrow \arg \max_c \bar{p}_c$ 
15:    Draw bounding box and overlay  $(\hat{y}, \bar{p}_{\hat{y}}, \bar{\mathbf{p}})$  on  $F$ 
16:  end for
17:  imshow( $F$ )
18:  if key = 'q' then break
19:  end if
20: end while
  
```

---

### 7.4 Computational Performance

On a mid-range GPU (e.g., NVIDIA GTX 1660):

- Haar Cascade:  $\approx 5$  ms per frame
- ResNet-18 forward pass:  $\approx 3$  ms per face
- Full pipeline:  $\geq 30$  fps for up to 4 simultaneous faces

On CPU only:

- ResNet-18 forward:  $\approx 25\text{--}40$  ms per face
- Sufficient for real-time demo (up to  $\approx 2$  faces at 20+ fps)

---

## 8 Grad-CAM Analysis

---

Grad-CAM (Section 4, Equations 24–25) was applied to 50 correctly classified test-set samples per class. The resulting attention regions are qualitatively consistent with known facial action units (FAUs):

Table 8: Qualitative Grad-CAM attention regions per emotion class.

Emotion	Primary Attention Region
angry	Inner brow area, nasal root
disgust	Upper lip, nose bridge
fear	Eyes (widened), brow region
happy	Mouth corners, cheeks
neutral	Diffuse (no dominant region)
sad	Mouth corners (downturned), inner brows
surprise	Eyes (wide), jaw

These patterns align with Ekman’s Facial Action Coding System (FACS) -an indication that the model has learned semantically meaningful representations rather than dataset artifacts. For `neutral`, the diffuse attention is expected: neutral is defined by the *absence* of activation rather than a specific configuration.

## 9 Limitations and Future Work

### 9.1 Limitations

**Closed-set assumption.** The model forces every input into one of seven classes. As Zhang *et al.* [1] demonstrate, this produces high-confidence mispredictions on compound or ambiguous expressions. The open-set framework (AUROC 0.909 on RAF-DB) could be integrated as future work.

**Illumination sensitivity.** FER2013’s wild-collected images include varied lighting, but the Haar Cascade face detector degrades significantly under very low light or strong lateral illumination, causing missed detections.

**Single-frame backbone.** The model processes each frame independently. Temporal models (3D CNNs, LSTM over frame sequences) could capture the dynamics of expression change, which is itself informative.

**Dataset biases.** FER2013 is known to have higher accuracy for Caucasian faces due to collection methodology. The model inherits this demographic bias.

### 9.2 Future Directions

- **Open-set extension:** Implement the noisy label detection framework of [1] to detect out-of-distribution expressions.
- **Multi-task learning:** Jointly predict emotion and facial action units to share feature representations.

- **Embedding-based FER:** Explore metric learning (triplet loss as in FaceNet [2]) for few-shot emotion recognition.
- **AffectNet training:** AffectNet contains 450,000+ images. Fine-tuning on it before FER2013 would provide a stronger facial expression prior.
- **MobileNet deployment:** Replace ResNet-18 with MobileNetV3 for edge-device deployment.

## 10 Conclusion

---

This project demonstrates that a carefully designed transfer learning pipeline can achieve near- or above-human-level performance on FER2013 while maintaining real-time inference from a standard webcam. The 15-experiment trajectory documents an iterative, evidence-driven design process that identifies the contribution of each component: transfer learning (+14% over scratch), staged unfreezing (+2%), grayscale adaptation (+1.7%), class weighting (+0.7%), label smoothing (+1.1%), and random erasing (+1.3%).

The mathematical formulation grounds every design decision in established theory: residual gradient flow (Equation 5), label smoothing distribution (Equation 14), Grad-CAM attribution (Equations 24–25), and temporal smoothing convexity (Equation 26).

The final system achieves weighted F1 of 0.71 on FER2013, produces interpretable Grad-CAM attention maps aligned with facial action unit theory, and runs stably at 30+ fps on a mid-range GPU.

## References

---

- [1] Y. Zhang, R. Su, W. Liu, and M. Wang, “Open-set facial expression recognition,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024. arXiv:2401.12507.
- [2] F. Schroff, D. Kalenichenko, and J. Philbin, “FaceNet: A unified embedding for face recognition and clustering,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2015, pp. 815–823. arXiv:1503.03832.
- [3] C. Dewi, R.-C. Chen, S.-H. Liu, and X. Jiang, “Real-time facial expression recognition: Advances, challenges, and future directions,” *Vietnam Journal of Computer Science*, 2024.
- [4] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778.

- [5] I. J. Goodfellow, D. Erhan, P. L. Carrier, A. Courville, M. Mirza, B. Hamner, W. Cukierski, Y. Tang, D. Thaler, D.-H. Lee, Y. Zhou, C. Ramaiah, F. Feng, R. Li, X. Wang, D. Melos, H. Wester, L. Bottou, and Y. Bengio, “Challenges in representation learning: A report on three machine learning contests,” in *Neural Networks*, 2013.
- [6] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-CAM: Visual explanations from deep networks via gradient-based localization,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2017, pp. 618–626.
- [7] P. Viola and M. Jones, “Rapid object detection using a boosted cascade of simple features,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2001.
- [8] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *International Conference on Learning Representations (ICLR)*, 2015. arXiv:1412.6980.
- [9] N. S. Keskar, D. Mudigere, J. Nocedal, M. Smelyanskiy, and P. T. P. Tang, “On large-batch training for deep learning: Generalization gap and sharp minima,” in *International Conference on Learning Representations (ICLR)*, 2017.